

# Kubernetes MultiNode cluster on VirtualBox

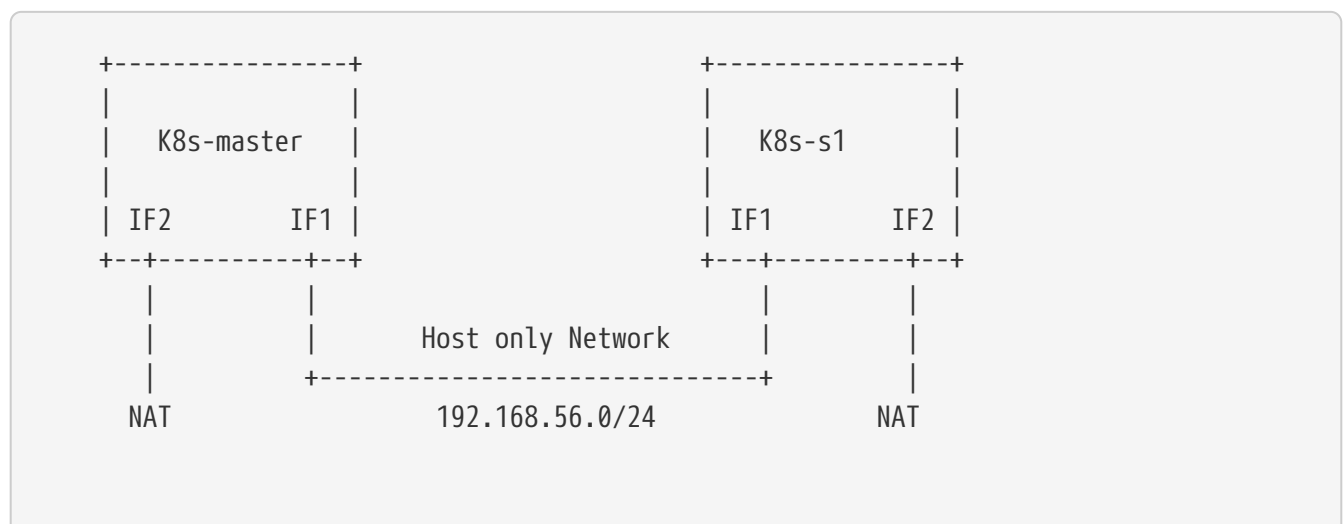
## Table of Contents

VirtualBox setup .....	1
VM hardware configuration.....	2
Prepare Servers.....	2
Setup network.....	2
Install docker.....	4
Install and configure kubernetes.....	5
Prepare server .....	5
Configure Master VM.....	8
Joining worker node to cluster.....	10
Install Helm v3 .....	11
Configure Kubernetes Dashboad (v2).....	11
Installing Metrics server.....	14
Installing Traefik Ingress Controller.....	15
Installing MetalLB.....	15
Validation.....	16
Troubleshoot.....	16

A step by step description of bringing up kubernetes cluster using VirtualBox 6.0 with ubuntu 18.04 server.

## VirtualBox setup

Create two VMs k8s-master and k8s-worker as shown in below diagram.



## VM hardware configuration.

- Kubernetes master/controller node hardware requirement

Component	Value (min)
Processor	2 core
RAM	4GB
Network Adapter-1	Host only network
Network Adapter-2	NAT

- kubernetes worker/slave node hardware requirement

Component	Value (min)
Processor	1 core
RAM	2GB
Network Adapter-1	Host only network
Network Adapter-2	NAT

## Prepare Servers

Execute the below step in both the VMs.

### Setup network.

- Shutdown the VMs.
- Modify VM network, with Adapter 1 for Host Only Network and Adapter 2 for NAT.
- Enable Nested Virtualization

```
$ VBoxManage modifyvm <vm-name> --nested-hw-virt on
```

*Example*

```
$ VBoxManage modifyvm k8s-master --nested-hw-virt on
$ VBoxManage modifyvm k8s-qoekwe --nested-hw-virt on
```



Nested Virtualization is supported for Intel i7 7Gen+ processors

- Start VM.

- Execute below command and note down the interfaces (generally for Virtual box VMs network interface will start with enp0s).

```
$ sudo ifconfig -a
```

- Setup static network.

Open file `/etc/netplan/00-installer-config.yaml` and add below configuration.

#### **k8s-master**

```
network:
  ethernets:
    enp0s3:
      addresses: [192.168.56.10/24]
      dhcp4: false
    enp0s8:
      addresses: []
      dhcp4: true
  version: 2
```

#### **k8s-worker**

```
network:
  ethernets:
    enp0s3:
      addresses: [192.168.56.12/24]
      dhcp4: false
    enp0s8:
      addresses: []
      dhcp4: true
  version: 2
```

- Update `/etc/hosts` such that command `hostname -i` outputs ip which is routable from other VM.

#### *example(master)*

```
192.168.56.10 k8s-master
192.168.56.12 k8s-worker
```

#### *example(worker)*

```
192.168.56.10 k8s-master
192.168.56.12 k8s-worker
```



This step is required since we've chosen host-only adapter and our default route is provided by NAT network, if we choose Bridge network this step can

be skipped.

- Reboot VMs.
- Update Packages.
  - Update apt sources.

```
$ sudo apt update
```

- Install packages

```
$ sudo apt upgrade
```

## Install docker

- Install utility packages

```
$ apt-get install -y \  
  apt-transport-https \  
  ca-certificates \  
  curl \  
  gnupg-agent \  
  software-properties-common
```

- Add GPG key

```
$ curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -
```

- Add docker repository to apt sources.

```
$ sudo add-apt-repository \  
  "deb [arch=amd64] https://download.docker.com/linux/ubuntu \  
  $(lsb_release -cs) \  
  stable"
```

- Update apt sources and install docker

```
$ sudo apt update ; sudo apt install -y docker-ce docker-ce-cli containerd.io
```

- Verify docker service is started.

```
$ sudo systemctl status docker
```

# Install and configure kubernetes.

## Prepare server

Execute the below steps on both the VMs.

1. Letting iptables see bridged traffic

```
cat <<EOF | sudo tee /etc/sysctl.d/98-kubernetes-cni.conf
net.bridge.bridge-nf-call-ip6tables = 1
net.bridge.bridge-nf-call-iptables = 1
EOF
```



Make sure that the `br_netfilter` module is loaded before this step. This can be done by running `lsmod | grep br_netfilter`. To load it explicitly call `sudo modprobe br_netfilter`.

2. Reload configuration

```
sudo sysctl --system
```

3. Add Kubernetes source gpg.

```
$ curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | sudo apt-key add -
```

4. Next add kubernetes repository

```
$ sudo apt-add-repository "deb http://apt.kubernetes.io/ kubernetes-xenial main"
```



Browse <http://apt.kubernetes.io/> and search for the corresponding ubuntu version, since there was no `kubernetes-bionic`, installing `kubernetes-xenial` in my case.

5. Install **kuber**enetes packages

```
$ sudo apt update; sudo apt install -y kubeadm kubelet kubectl
```

6. Prevent automatic update of packages

```
$ sudo apt-mark hold kubelet kubeadm kubect1
```

7. Turn off swap

```
$ sudo swapoff -a
```

8. Turn off swap permanently in `/etc/fstab` by executing below command.

```
$ sudo sed -i '/swap/s/^\(.*\)$/#\1/g' /etc/fstab
```

9. Add private ip to kubelet. Create file `/etc/default/kubelet` and following content in the file.

```
KUBELET_EXTRA_ARGS=--node-ip=192.168.56.10
```



Change ip address as per master and worker node ips

10. Configure docker `daemon.json`.

```
cat <<EOF | sudo tee /etc/docker/daemon.json
{
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
EOF
```

11. Install ipvs, we shall be using ipvs for LB.

```
sudo apt install -y ipvsadm
```

12. configure ipvsadm

```
cat <<EOF | sudo tee /etc/sysctl.d/98-ipvsadm.conf
net.ipv4.ip_forward=1

net.ipv4.conf.all.arp_ignore=1
net.ipv4.conf.all.arp_announce=2
EOF
```

### 13. update ipvsadm configuration

```
cat <<EOF | sudo tee /etc/default/ipvsadm
# ipvsadm

# if you want to start ipvsadm on boot set this to true
AUTO="true"

# daemon method (none|master|backup)
DAEMON="master"

# use interface (eth0,eth1...)
IFACE="enp0s3"

# syncid to use
# (0 means no filtering of syncids happen, that is the default)
# SYNCID="0"
EOF
```



Change the network interface accordingly

### 14. Reload configuration

```
sudo systemctl --system
```

### 15. Load modules

```
sudo modprobe -- ip_vs
sudo modprobe -- ip_vs_rr
sudo modprobe -- ip_vs_wrr
sudo modprobe -- ip_vs_sh
sudo modprobe -- ip_vs_sed
sudo modprobe -- ip_vs_lblc

sudo lsmod | grep -e ip_vs -e nf_conntrack
```

### 16. Reload configuration and restart **docker** and **kubelet** service.

```
sudo systemctl daemon-reload
sudo systemctl restart docker
sudo systemctl restart kubelet
sudo systemctl restart ipvsadm
```

### 17. **(Optional)**Reboot VMs.

# Configure Master VM.

1. Create `kubeadm-config.yaml` file that consists of cluster configuration.

```
$ cat <<EOF | tee kubeadm-config.yaml
apiVersion: kubeadm.k8s.io/v1beta2
bootstrapTokens:
- groups:
  - system:bootstrappers:kubeadm:default-node-token
  ttl: 24h0m0s
  usages:
  - signing
  - authentication
kind: InitConfiguration
localAPIEndpoint:
  advertiseAddress: 192.168.56.10
  bindPort: 6443
---
apiServer:
  timeoutForControlPlane: 4m0s
apiVersion: kubeadm.k8s.io/v1beta2
certificatesDir: /etc/kubernetes/pki
clusterName: kubernetes
#controlPlaneEndpoint: 192.168.1.6:6443
controllerManager: {}
dns:
  type: CoreDNS
etcd:
  local:
    dataDir: /var/lib/etcd
imageRepository: k8s.gcr.io
kind: ClusterConfiguration
#kubernetesVersion: v1.18.0
networking:
  dnsDomain: cluster.local
  serviceSubnet: 10.96.0.0/12
  podSubnet: 10.244.0.0/16
scheduler: {}
FeatureGates:
  ServiceTopology: true
  TopologyManager: true
---
apiVersion: kubeproxy.config.k8s.io/v1alpha1
bindAddress: 0.0.0.0
ipvs:
  scheduler: "lbic"
  strictARP: true
kind: KubeProxyConfiguration
mode: "ipvs"
```





If decide to change pod network, remember to change the pod network in `flannel.yaml` while setting up pod network in below step.



You can geneate the above configuration file using below command

```
$ sudo kubeadm config print init-defaults --component-configs
KubeProxyConfiguration
```

## 2. Bring up cluster now.

```
$ sudo kubeadm init --config kubeadm-config.yaml
```

## 3. As output of above command shows to execute below command, run below command to update auth details for `kubectl` command.

```
$ mkdir -p $HOME/.kube
$ sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
$ sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

## 4. Note down (copy to notepad) the `kubeadm` join command which needs to be executed on worker nodes to join to the cluster.

## 5. Install pod network.

- a. If you are using bridge network which has default route then you can directly install flannel using below command. However if you had given host-only adapter IP address for API server. then skip this step and follow next steps.

```
$ kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yaml
```

### b. Download flannel.yaml

```
$ wget https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yaml
```

- c. Open `kube-flannel.yaml` file search for daemonset kind and add additional arugment `--iface=<host-only-adaper-interface>`, note there are multiple daemonset definitions, either update all of them or execute `kubectl get nodes --show-labels` and get the arch and update corresponding arch daemonset only.

```
containers:
- name: kube-flannel
  image: quay.io/coreos/flannel:v0.11.0-amd64
```

```

command:
- /opt/bin/flanneld
args:
- --ip-masq
- --kube-subnet-mgr
- --iface=enp0s3 ①
resources:
  requests:
    cpu: "100m"
    memory: "50Mi"
  limits:

```

① Added my host-only adapter here

- d. if you had changed pod network while executing `kubeadm init`, replace `10.244.0.0/16` which chosen pod network above in `net-conf.json` section.
- e. Now Create pod network by applying updated yaml file

```
$ kubectl apply -f kube-flannel.yml
```

## 6. Verify all necessary pods are started

```
$ kubectl get pods -A
```

output:

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	coredns-86c58d9df4-c68gd	1/1	Running	0	6m41s
kube-system	coredns-86c58d9df4-q5bht	1/1	Running	0	6m41s
kube-system	etcd-k8s-master	1/1	Running	0	6m6s
kube-system	kube-apiserver-k8s-master	1/1	Running	0	5m59s
kube-system	kube-controller-manager-k8s-master	1/1	Running	0	5m56s
kube-system	kube-flannel-ds-amd64-stb29	1/1	Running	0	49s
kube-system	kube-proxy-882ms	1/1	Running	0	6m41s
kube-system	kube-scheduler-k8s-master	1/1	Running	0	5m54s

## 7. We can check if master node is ready or not by executing below command.

```
$ kubectl get nodes
```

# Joining worker node to cluster.

1. Now go to worker node and execute the join command previously saved when you were executing `kubeadm` on master.

```
$ sudo kubeadm join 192.168.56.10:6443 --token t0j1zi.v5lojsnpjh9r0rbn \
--discovery-token-ca-cert-hash sha256:40b1142d9002003ab5b085776b8b8cba4a41ceaafab06429c49eadc2b2939fa
```



The above command is sample, the values are dynamically generated.

2. Now go back to master and execute the below command, you should be able to see slave node added.

```
$ kubectl get nodes
```

*output:*

NAME	STATUS	ROLES	AGE	VERSION
k8s-master	Ready	master	56m	v1.18.4
k8s-worker	Ready	<none>	2m49s	v1.18.4

## Install Helm v3

Helm package manager consists of helm client [ Helm v3 deprecated need for server side component(tiller) which was running as pod on k8s]. We will install helm client on our master server itself.

### Install helm client.

1. Install helm client using script on master server using below set of commands.

```
$ curl -fsSL -o get_helm.sh
https://raw.githubusercontent.com/helm/helm/master/scripts/get-helm-3
$ chmod 700 get_helm.sh
$ ./get_helm.sh
```

2. Verify Helm is running

```
$ helm list
```

*output*

NAME	NAMESPACE	REVISION	UPDATED	STATUS	CHART	APP VERSION
------	-----------	----------	---------	--------	-------	-------------

## Configure Kubernetes Dashboad (v2).

## Install Dashboard

1. Add dashboard repository.

```
$ helm repo add kubernetes-dashboard https://kubernetes.github.io/dashboard/
```

2. Update dashboard repository.

```
$ helm repo update
```

3. Install dashboard.

```
helm install kubernetes-dashboard \
  kubernetes-dashboard/kubernetes-dashboard \
  --namespace kube-system \
  --set fullnameOverride=kubernetes-dashboard \
  --set serviceAccount.name=admin-user \
  --set metricsScraper.enabled=true \
  --set service.type=NodePort
```



By default dashboard is not recommended to be accessed from outside the VM, if you are using ubuntu desktop you can run below command and access the dashboard using proxy at url <http://localhost:8001/api/v1/namespaces/kube-system/services/https:kubernetes-dashboard:/proxy/> however we will be choosing **NodePort** to access dashboard using host vm's IP as we set **service.type=NodePort**.

+

```
$ kubectl proxy
```

4. Wait till dashboard pod is running.

```
$ kubectl get pods -n kube-system -w
```

output:

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	coredns-86c58d9df4-c68gd	1/1	Running	0	11m
kube-system	coredns-86c58d9df4-q5bht	1/1	Running	0	11m
kube-system	etcd-k8s-master	1/1	Running	0	10m
kube-system	kube-apiserver-k8s-master	1/1	Running	0	10m
kube-system	kube-controller-manager-k8s-master	1/1	Running	0	10m
kube-system	kube-flannel-ds-amd64-stb29	1/1	Running	0	5m18s
kube-system	kube-proxy-882ms	1/1	Running	0	11m



Copy the token and paste it into token field in the URL to Dashboard and login to dashboard.



## 1. Add stable repository

## 2. Update helm repository

3. Execute below helm command to install metrics server.

4. Wait for pod to come up after that you can execute top command to get resource usage.

5. Alternatively instead of using helm, download latest released component.yaml from metrics-server [github release page](#) and then add following two args to metrics-server container[in args section after below secure-port=4443]. and finally install using `kubectl apply -f component.yaml`

14

# Installing Traefik Ingress Controller.

1. Add traefik repository

```
$ helm repo add traefik https://containous.github.io/traefik-helm-chart
```

2. Update helm repository

```
$ helm repo update
```

3. Execute below helm command to install traefik.

```
$ helm install traefik traefik/traefik --namespace kube-system
```

# Installing MetalLB.

1. Execute following commands, This will deploy MetalLB to your cluster, under the metallb-system namespace

```
$ kubectl apply -f
https://raw.githubusercontent.com/metallb/metallb/v0.9.3/manifests/namespace.yaml
$ kubectl apply -f
https://raw.githubusercontent.com/metallb/metallb/v0.9.3/manifests/metallb.yaml
$ kubectl create secret generic -n metallb-system memberlist --from
-literal=secretkey="$(openssl rand -base64 128)"
```

2. Create L2 configuration into config-map file.

```
cat <<EOF | tee metallb-l2-config.yaml
apiVersion: v1
kind: ConfigMap
metadata:
  namespace: metallb-system
  name: config
data:
  config: |
    address-pools:
    - name: default
      protocol: layer2
      addresses:
      - 192.168.56.240-192.168.56.250
EOF
```

3. Apply the configuration

```
kubectl apply -f metallb-l2-config.yaml
```

4. Now if you can list service and notice that **external ip** field has ip address set from above pool. This can be noticed for traefik service.

```
kubectl get svc -n kube-system
```

## Validation.

You can deploy a sample angular application uploaded [here](#)

## Troubleshoot.

<https://kubernetes.io/docs/tasks/administer-cluster/dns-debugging-resolution/> - DNS not resolving

<https://kubernetes.io/docs/setup/independent/troubleshooting-kubeadm/> - kubeadm

<https://github.com/kubernetes/kubernetes/tree/master/pkg/proxy/ipvs> - ipvs

## References

<https://github.com/helm/charts/tree/master/stable/traefik> - traefik

<https://github.com/helm/charts/tree/master/stable/nginx-ingress> - nginx

<https://kubernetes.io/docs/concepts/services-networking/ingress/> - ingress controller

<https://metallb.universe.tf> - metallb